

Project Mthreads

Due: March 3, 2017

Contents

1	Introduction	1
2	Background	1
3	The Assignment	2
3.1	Overview	3
3.2	Assumptions	4
3.3	Topics from Uthreads	4
3.4	The Reaper 2.0	4
4	Compiling and Testing	5
4.1	Debugging	5
4.2	Test Code	5
4.3	Working From Home	5
5	Handing In	6

1 Introduction

In this project you will build on the TA solution for the *uthreads* assignment to make it multiprocessor-safe. In *uthreads*, you implemented a threading library with an N-to-1 (N uthreads multiplexed on 1 “kernel” thread) threading model. In *mthreads*, you will be implementing a library with an M-to-N (M uthreads multiplexed on N “kernel” threads) threading model. You can think of **pthreads** as kernel threads. Having completed this assignment, your threading library will support user-level threads running *in parallel* on many processors! You can copy the stencil code out of the `/course/cs167/asgn/mthreads_stencil` directory.

2 Background

M-to-N, also known as *Hybrid Threading*, is a threading model in which user-level application threads are multiplexed on a pool of kernel-level threads. Such kernel-level threads can be referred to as “virtual processors”, since these threads may or may not actually be running in parallel – N kernel threads may themselves share a single CPU.

In *uthreads*, each user-level thread had a context that the kernel thread, which was actually the main thread of the program, switched into when the associated thread was scheduled to run. In *mthreads*, the story is basically the same only that instead of having just one main thread, we have a number of **pthreads**, each of which represents a kernel thread in an OS. These **pthreads** are referred to in the stencil as a lightweight-processes (LWP). LWP's execute one runnable **uthread** at a time.

M-to-N threading is considerably more complicated to implement than N-to-1 threading because, in addition to masking preemption in the right places, the model must synchronize kernel-thread access to data structures including, but not limited to, the run queue, user-level mutexes/condition variables, and user-level thread structures. In fact, this is most of the work you will be doing.

When implementing this assignment it is also important to keep in mind that a user-level thread may be scheduled on any of a number of kernel threads and processors during its lifetime. This means that in a routine like **uthread_switch**, we need to ensure that data that is necessary before and after a user thread is scheduled is not optimized away into a CPU register by the compiler.

Finally, you will find that there are several instances (documented in the stencil) where you will have to perform some operation on a user-thread member *after* the thread in question has switched off the processor. Reasons for doing this are described in the stencil, but you are strongly encouraged to read chapter 5 of the textbook, which discusses various multiprocessor issues like this. Particular egregious errors include two processors executing the same thread, and a currently-executing thread's stack being freed while some user-level thread is executing on that stack.

3 The Assignment

mthreads will make use of a lot of code that should already be familiar to you from the *uthreads* assignment. The logic for creating threads, joining and detaching threads, scheduling, and dealing with mutexes and condition variables is the same, only now you will find that the concerned structures have been augmented with a **pthread_mutex_t**, which should be used in the appropriate contexts to synchronize various LWP's contending for access.

You only need to worry about modifying functions with TODOs. Each of these also contains a call to our **NOT_YET_IMPLEMENTED** macro. To see a list of all the functions you have yet to modify/implement, run **make nyi** in the project working directory.

The majority the work in the assignment will go into implementing **uthread_switch**, which should be significantly different from the one you implemented in *uthreads*, and **lwp_switch**, which is where LWP's execute independent of any **uthread** context, waiting for a runnable **uthread** to become available. Here's a list of functions you'll need to touch while working on this assignment:

```
// uthread.c
uthread_exit(int status);
uthread_join(uthread_id_t uid, int *return_value);
uthread_detach(uthread_id_t uid);
make_reapable(uthread_t *uth);

// uthread_sched.c
uthread_yield(void);
```

```

uthread_block(void);
uthread_wake(uthread_t *uth);
uthread_setprio(uthread_id_t id, int prio);
uthread_startonrunq(uthread_id_t id, int prio);
uthread_switch(utqueue_t *q, uthread_t *thr, pthread_mutex_t *m);
lwp_switch(void);
lwp_park(void);
uthread_runq_enqueue(uthread *thr);
uthread_runq_requeue(uthread *thr);

// uthread_mtx.c
uthread_mutex_lock(uthread_mtx_t *mtx);
uthread_mutex_trylock(uthread_mtx_t *mtx);
uthread_mutex_unlock(uthread_mtx_t *mtx);

// uthread_cond.c
uthread_cond_wait(uthread_cond_t *cond, uthread_mtx_t *mtx);
uthread_cond_broadcast(uthread_cond_t *cond);
uthread_cond_signal(uthread_cond_t *cond);

```

It may help to implement `uthread_switch` and `lwp_switch` before proceeding to the other functions even though they require more work because a) these routines interact with each other in a somewhat complicated way and b) they make the majority of the heavy-weight scheduling decisions that involve switching between the context of a `uthread` and the context of an LWP. The understanding you gain from implementing these two functions should make editing the remaining ones much easier.

3.1 Overview

Many of the functions mentioned above (particularly `lwp_switch`) have extensive comments in the source code which explain what is expected of you, but to save you some time, we will give you a brief summary of how the system works as a whole.

The first thing that any executable that uses your threads package should call is `uthread_start`. This should be called exactly once and is responsible for setting up global data structures such as the `uthreads` array (as it did in *uthreads*). Note that it, in turn, calls the function to be executed by the first user thread of the program.

The major differences between the setup done by `uthread_start` and `uthread_init` (from *uthreads*) lies in `create_first_thr`. This function is responsible for setting up the main thread and the reaper thread, and placing them both on the run queue. It also sets up additional LWPs by calling `uthread_start_lwp` which spawns a pthread that runs `lwp_start`. It is important to note that, at this point, each of these two new LWPs has its own `ut_curthr` and `curlwp` variables. This can be seen from the declarations near the top of `uthread.c`:

```

// uthread.c
__thread uthread_t *ut_curthr = 0;
__thread lwp_t    *curlwp;

```

The `__thread` storage class keyword is an extension to GCC that can be used alone or with the `extern` or `static` qualifiers for a global, file-scoped static, or function-scoped static variable to ensure that each thread is allocated its own private copy of that variable. Addresses of such variables can be used by other threads, though you shouldn't need to worry about that for this assignment. This type of storage is known as thread-local storage (TLS). GCC implements it for POSIX threads (which you use as your LWPs; not for user threads. Note that TLS is not async-signal safe. This means that if as signal handler can access a TLS item, you must make sure that when threads access the item that signals are masked.

What this means for us is that each LWP stores its own context *in addition to* the context of the `uthread` that it is currently running. Thus we can expect that `ut_curthr` will change frequently throughout the LWP's lifetime (in `lwp_switch`) as it multiplexes various user-level threads, whereas `curlwp` will remain constant so that various `uthreads` can switch back into the context of their invoking LWPs if they need to block or yield. This is exactly what the interaction between `uthread_switch` and `lwp_switch` facilitates. A `uthread` invoking `uthread_switch` uses its thread-local `curlwp` to jump back into `lwp_switch`, whereas an LWP users the `ut_ctx` member of a runnable `uthread` to switch into its context and begin running.

By the end of `uthread_start` all of the LWPs are ready to begin multiplexing `uthreads`. The remainder of the assignment consists of ensuring that the `uthread` API (that you implemented in the previous assignment) is thread-safe *and* multi-processor safe. The `mthreads` library still supports user-level preemption, however your TAs have provided you with a solution to `uthreads` that deals with this appropriately in most of the `uthreads` API you will be working with.

3.2 Assumptions

In this assignment there will be multiple threads running at a time, and on multiple processors. Thus, make sure you must guard data structures appropriately. Be sure to review `uthread.h`, which contains definitions for `uthread_t` and `lwp_t`, as well as other struct definitions for hints as to which data structure accesses need to be synchronized. Importantly, you **cannot** make the assumption that each `uthread` is scheduled on exactly one LWP at a time. This is something that you must ensure with proper synchronization!

The stencil as provided will not work because `uthreads` aren't synchronized, so running the stencil as given will not give you any information. You may, however, assume that any functions that do not have TODOs will work properly once you make the required functions (described in the **Assignment** section) thread-safe and multiprocessor-safe. You will not be held accountable for potential problem with the stencil, but you **must** modify your API to work with the stencil. If you write alternate implementation that doesn't work with the given stencil, you will receive no credit.

3.3 Topics from Uthreads

Understanding the concepts from `uthreads` will be very important in `mthreads`. You may want to reread the `uthreads` assignment, including the sections Swapping Contexts, Time Slicing and Preemption, Dangers of Preemption, and the Reaper.

3.4 The Reaper 2.0

Conceptually, the reaper in *mthreads* serves the same purpose it did in *uthreads*. However, because the reaper can now run in parallel with other *uthreads*, we must be sure that it doesn't clean up resources associated with a thread that has just called `uthread_exit` before that thread has switched away from its context and is no longer calling functions on that context's stack. For this reason, before the reaper calls `uthread_destroy`, it must first attempt to lock the exited thread's *ut_pmut* mutex to be sure that it is no longer using its context. It is up to you to implement your *uthread* API so that the reaper can safely do its job.

4 Compiling and Testing

Currently, it is not possible to take any (already compiled) program and have it use *mthreads* instead of *pthreads*. In order to use *mthreads*, you will need to add all of *mthreads*'s files to its project, modify it to use the *mthreads* API functions and recompile it.

4.1 Debugging

As always, use of `gdb` will make your life much easier when trying to get this assignment up and running. However, `gdb` can sometimes get confused in multithreaded programs, and may have trouble printing stack traces. If that happens, don't worry; it doesn't mean your code is broken. Also, since *uthreads* is built as a library, `gdb` won't find the symbols in it right away, so tell it to wait for the "future shared library to load".

4.2 Test Code

We can't stress enough how important test code is in an assignment like this. Without proper test code, finding bugs will be next to impossible. Make sure to test all sorts of situations with lots of threads at different priority levels. The `Makefile` included with the assignment will compile a reasonably thorough test program (`preemptive_test`) which uses the *uthreads* functions. If it runs and exits cleanly, most of your basic functionality is working, but we encourage you to write your own tests as well.

Judicious use of `assert()` will help you both understand your threads package and debug it. This is your first real systems-level coding project, and it is highly recommended that you assert a general sane state of the system whenever you enter a function. Thinking about what a "sane state" means should lead to a greater understanding of what is happening at any given time and what could go wrong.

A final warning: `printf()` is NOT async-signal-safe. If you are going to write a program to test your *mthreads*, consider using a combination of `sprintf()` and `write()` like is done in the test program the TAs provide for you. In general for this assignment, it is not safe to call a function which isn't async-signal-safe unless preemption has been masked, which means "user" code should only call async-signal-safe functions. Think about why this is the case.

4.3 Working From Home

Note that *mthreads* will not work on OSX, so if you work from home you may not be able to run *mthreads* natively.

5 Handing In

To hand in your finished assignment, please run this command while in the directory containing your code: `/course/cs167/bin/cs167_handin mthreads`.